

Highlights

Typed Issue Specifications from App Reviews: What Scales, and a Downstream Benefit That Does Not Replicate

Fabiha Jalal, Sadia Tasnim Dhruba, Tazrian Zaman Tushti, Ajwad Abrar, Md Kamrul Hasan, Hasan Mahmud

- A typed intermediate representation routes app reviews into five issue templates.
- A verified-anchor procedure lifts Cohen's kappa from 0.163 to 0.592 at corpus scale.
- LLM annotation noise concentrates on omitted classes, not on uniform per-item error.
- A clean same-model re-run finds no downstream reply-quality gain from the typed spec.
- Template-fill measures schema conformance, not the quality of the text in the fields.

Typed Issue Specifications from App Reviews: What Scales, and a Downstream Benefit That Does Not Replicate

Fabiha Jalal^{a,*}, Sadia Tasnim Dhruba^a, Tazrian Zaman Tushti^a, Ajwad Abrar^b, Md Kamrul Hasan^b, Hasan Mahmud^b

^a*Department of Computer Science, Islamic University of Technology (IUT), Gazipur, Bangladesh*

^b*Systems and Software Lab (SSL), Department of Computer Science and Engineering, Islamic University of Technology (IUT), Gazipur, Bangladesh*

Abstract

Context. Defect trackers require typed fields (issue type, steps to reproduce, severity, affected component) and cannot consume the free text users write in app reviews. Review-mining research stops short of that boundary, predicting labels or writing replies without producing the record a tracker can ingest.

Objective. Can a typed intermediate representation (IR) of a review be produced at corpus scale from noisy labels, and does supplying it to a response generator improve the reply the user receives?

Method. IssueSpec routes each review by issue class into one of five standards-body templates and repairs the noisy labels the routing depends on against a human-verified anchor. We study 215,583 reviews from 58 Android apps with a 5,230-review anchor, a 490-review three-rater gold standard, seven ablations and four language models. The downstream evaluation was run twice, the second time holding the generator fixed.

Results. The verified-anchor procedure lifts Cohen’s κ against the human gold from 0.163 to 0.592, holding at that level on held-out reviews. Type routing reaches 0.96 template-fill against 0.69 without it, a schema-

*Corresponding author.

Email addresses: fabihajalal@iut-dhaka.edu (Fabiha Jalal), sadiatasnim@iut-dhaka.edu (Sadia Tasnim Dhruba), tazrian@iut-dhaka.edu (Tazrian Zaman Tushti), ajwadabrar@iut-dhaka.edu (Ajwad Abrar), hasank@iut-dhaka.edu (Md Kamrul Hasan), hasan@iut-dhaka.edu (Hasan Mahmud)

conformance measure, not a quality one. The downstream benefit does not hold. An earlier round measured +2.35 Likert points, but its two arms came from different template composers rather than one model; corrected, the gain is +0.04 ($p = 0.54$), with the three raters disagreeing in direction. Three further components fail their own tests: the knowledge-graph layer loses to count-matched flat clustering, our faithfulness scorer shows no rank correlation with human judgement, and the constrained alignment layer never demonstrates its intended trade-off.

Conclusions. Typed issue specifications are producible at corpus scale, and that is the contribution we defend. Pipelines of this shape should be evaluated on what the intermediate representation produces rather than on its effect downstream, and a same-model comparison should be the default before any such effect is claimed.

Keywords: App review mining, Issue report generation, Intermediate representation, Knowledge graph, Retrieval-augmented generation, Empirical software engineering

1. Introduction

App stores are where users talk back to developers, and the global stream runs into millions of¹ reviews a day [1, 2, 3]. Those reviews carry bug reports, feature requests, performance complaints, and regression notes as messy free-form text, usually without the details required to reproduce a bug or scope a feature. Defect trackers such as JIRA or GitHub Issues cannot consume raw text; they need structured fields such as issue type, steps to reproduce, severity, and affected component [4, 5], and filling those by hand does not scale. Years of work on automated review pipelines has helped with classification, summarisation, and response generation, but no system produces the structured report a developer can paste into a tracker. Dąbrowski et al.’s systematic review of 182 papers [3] surveys this literature; it does not itself name this gap, and we state the gap as our own reading of the work it covers. Section 2 treats each strand in turn: classifiers stop at the label, response

¹Abbreviations: IR, intermediate representation; KG, knowledge graph; RAG, retrieval-augmented generation; RLHF, reinforcement learning from human feedback; CMDP, constrained Markov decision process; HITL, human-in-the-loop; IRB, institutional review board.

generators never write the defect out in structured form, agentic software-engineering tools assume a well-formed issue as the agent’s starting state, knowledge-graph methods emit predictions and summaries a tracker cannot ingest, and constrained RLHF targets general LLM safety rather than the domain-specific constraint family of developer relations.

This work pursues three aims, stated as research questions.

RQ1 (Translation Quality). Can a typed IR, produced by a three-layer framework (knowledge-graph construction, hierarchical clustering, standardised schema mapping), reach the substantive completeness of human-written GitHub issues?

RQ2 (Downstream Load-Bearing). Does conditioning a response generator on the validated IssueSpec measurably improve human-rated response quality over RAG alone? Section 5.2 answers no, correcting our own earlier claim.

RQ3 (Dual-Objective Alignment). Can a CMDP-grounded RLHF loop through KTO, DPO, and Constrained PPO jointly optimise response helpfulness and policy compliance for developer relations? Section 5.3 answers no at the scale we can train, and reports what that rules out.

We propose **IssueSpec**: a typed Intermediate Representation (IR) that converts each noisy review into one of five standard report types, and a CMDP-grounded response generator that consumes it under quality and compliance constraints (Figure 1).

What is new, and what is assembled. Several components are existing technique used as intended, and we claim none of them: RoBERTa classification, confident learning [6], Sentence-BERT with UMAP and HDBSCAN, PageRank, retrieval-augmented generation, and Altman’s CMDP formulation [7]. Three things are new. *First*, the typed intermediate representation: a review is routed by predicted issue class into one of five standards-body templates, so the output is a typed record with per-type obligations rather than a label or a summary, which no prior review-mining system emits. *Second*, the verified-anchor procedure, which makes that representation producible from noisy LLM labels at corpus scale, and which works because the labeller omits two of its seven classes entirely rather than because its errors concentrate on rare ones. *Third*, a correction and the check that produced it: we previously reported that the typed intermediate carries downstream response quality, and Section 5.2 withdraws that, because the comparison behind it varied the text generator as well as the specification, and because the arithmetic of the two rating rounds turned out to distinguish independent human ratings from

ones derived by perturbing another rater’s. We release both rounds and that check.

The central claim, stated once. The answer to which of five stages the paper is about is Stage 3 and the object it produces, a *typed, type-routed intermediate representation* of an app review, together with the verified-anchor procedure that makes it producible at corpus scale from noisy labels; everything else feeds that object or consumes it. Removing type routing drops strict template-fill from 0.96 to 0.69 (Section 5.1). We had also claimed a downstream payoff, and Section 5.2 reports why we withdrew it: with the generator held fixed, supplying the IR changes response quality by +0.04 Likert points ($p = 0.54$), against the +2.35 a confounded design gave. Three further components fail their own tests, the knowledge-graph layer against count-matched flat clustering, the SPECCOV faithfulness scorer against human judgement, and the CMDP alignment layer against the trade-off it was built to expose. Alongside these we release a reusable check on rating provenance, described in Section 5.2. A reader who wants the claim and not the machinery can read Sections 5.1, 5.2, and 6.1 alone. Table 1 is the single map from component to claim to evidence to verdict, and replaces the several overlapping summaries an earlier version carried.

To help others replicate or extend this work, we release the code, the annotation instruments, both rating rounds, and a verification script that recomputes the numbers in Section 5 from the released data. The data-availability statement gives the archived DOI.

The remainder of this paper is organised as follows. Section 2 positions the work against review mining, response generation, knowledge graphs, and constrained alignment. Section 3 describes the pipeline and the typed representation it produces. Section 4 gives the study design: the research questions, the corpora, the annotation protocol, the metrics, and the analysis procedure. Section 5 answers each research question in turn. Section 6 discusses what survives and what does not, Section 7 addresses threats to validity, and Section 8 concludes.

2. Related Work

We organise related work by the five gaps of Section 1, naming what each field does, what assumption breaks for review-to-issue translation, and how IssueSpec differs.

component	what it must do	evidence, and the ablation that tests it	verdict
Stage 1 in-take, verified anchor	repair a corpus label wrong 25% of the time that omits two classes entirely, because Stage 3 routes on the type	κ 0.163 $\rightarrow$ 0.592 against the three-rater gold, 0.616 held out (replay-dependent); macro F1 0.22 $\rightarrow$ 0.65 with the two omitted classes recovered (§5.1)	justified
Stage 2 knowledge graph	group and rank 215,583 reviews so a developer has an order to work in	count-matched flat clustering beats the KG on all three intrinsic metrics (A1b) and on purity (A2); top-5 PageRank aspects cover 55% of reviews	prioritisation yes, cluster quality no
Stage 3 typed IR	emit typed tracker fields rather than prose	strict fill 0.96 vs 0.69 without routing, i.e. schema conformance and not quality (A3, Table 3); 5-dim rubric 3.94/5 on all 100 specs, LLM judge (Table S2 in the supplementary material)	justified
Stage 4 spec-aware RAG	supply house voice and precedent on top of the defect	retrieval moves automatic metrics < 1% (A5); conditioning on the IR moves quality +0.04, $p = 0.54$ (A6, §5.2), overturning the +2.35 we reported	withdrawn / null
Stage 5 CMDP RLHF	hold compliance as a constraint rather than a reward term	constraint never binds at distil-GPT2 scale; with a policy able to violate it, binds on 13/90 steps and λ rises to 0.65, with no measured tension ($\rho = -0.12$, $p = 0.25$), on a detector that scores hand-written violations as compliant (§5.3)	partial testbed
SPEC-COV scorer	rank specs the way a human judge of faithfulness does	no rank correlation with human faithfulness ($\rho = 0.15$, $p = 0.26$); orders generators backwards (§5.1)	withdrawn
rating-provenance check	separate independent ratings from derived ones	the discarded round shows a single deviation magnitude and the reported round four (§5.2)	released

Table 1: Component-to-claim-to-evidence map, with the ablation that tests each component and the verdict it returns. “Removing this hurts” is the only outcome that justifies a component; several do not clear that bar and we report them as findings rather than as caveats. Ablation identifiers (A1–A7) are defined in Section 5.5; released artifacts are listed in Section 6.3.




Figure 1: The IssueSpec five-stage pipeline. Reviews flow left to right through intake and classification (Stage 1), three-layer knowledge-graph clustering (Stage 2), LLM-agent IR translation (Stage 3), spec-aware RAG (Stage 4), and CMDP-grounded RLHF (Stage 5), each emitting a typed intermediate. Yellow checkpoints mark human-in-the-loop review under selective prediction [8, 9]; the red loop propagates Stage 5 corrections upstream.

App review classification and mining. Maalej and Nabil [1] established a four-class taxonomy (bug report, feature request, user experience, rating). The seven-class label set we use is our own extension, adding **performance**, **usability**, and **compatibility** and renaming the rating class **praise**; we are not aware of an established seven-class app-review taxonomy, so the extension should be read as ours rather than as canonical. Subsequent systems [10, 11, 12, 13, 14, 15] refined accuracy and added aspect-level granularity, and Dąbrowski et al. [3] review 182 papers from 2012 to 2020, reporting median precision above 81% and median recall around 83%. An earlier version of this paper attributed to that survey a claim it does not make, that label quality rather than model capacity sets an F1 ceiling of 0.75 to 0.85: the survey reports precision and recall rather than F1 and attributes limitations to evaluation-dataset size, so the label-quality argument is ours and Section 5.1 is evidence for it rather than corroboration of theirs. *What breaks:*

none of these systems emits a typed developer-routable artifact downstream of the label.

Response generation and RAG. RRGGen [2] and CoRe [16] pioneered end-to-end neural response generation on (review, response) pairs without an explicit structured intermediate; vanilla RAG [17] retrieves past responses in one pass and agentic variants [18, 19] add critique-revise loops. *What breaks:* none of these names the underlying defect, and our $n = 10$ feasibility study (the supplementary material) shows that agentic iteration without a typed intermediate adds only a small gain on a rule-based scorer ($0.58 \rightarrow 0.70$). We cannot claim the IssueSpec does better: once the generator is held fixed, supplying it changes human-rated quality by $+0.04$ ($p = 0.54$, Section 5.2), so the contribution we defend is the specification itself rather than its effect on the reply.

Agentic software engineering. SWE-Agent [20], RepairAgent [21], and HyperAgent [22] consume well-structured GitHub issues to produce code patches, and the SWE-bench protocol [23] makes the assumption explicit by pairing each instance with an existing issue. *What breaks:* the specification is the input contract for all of them, and constructing one from noisy user-facing text is out of scope. IssueSpec is precisely the artifact these systems are designed to consume.

Knowledge graphs for app reviews. ReviewGraph [24] learns KG embeddings for rating prediction, KGCPN [25] produces extractive summaries, and Keertipati et al. [26] apply centrality for coarse feature prioritisation. *What breaks:* none maps KG nodes to a typed, developer-actionable schema. Our three-layer KG pipeline is, to our knowledge, the first that emits a typed IR for review-to-issue translation.

RLHF and constrained optimisation. DPO [27], KTO [28], and PPO [29] compress quality and compliance into one scalar reward; Altman’s CMDP framework [7] separates them, and Safe-RLHF [30] applies CMDP to general LLM safety with a domain-agnostic compliance set. *What breaks:* no prior work applies CMDP to dev-rel response generation, where the compliance set is domain-specific (promise safety, information leakage, tone, legal safety). Our Stage 5 formulation instantiates these four constraints inside the Lagrangian and releases a proof-of-concept training loop.

Across these categories, existing systems occupy at most two of our five pipeline stages and none produces a typed issue record. Composing all five with a standardised issue schema, three HITL checkpoints, KG-centrality prioritisation, and a CMDP-grounded RLHF formulation is the structural

contribution.

3. The IssueSpec Pipeline

The pipeline (Fig. 1) is a five-stage information-extraction cascade [31, 32]. Stages 1–3 produce the typed IR from raw reviews, Stage 4 uses it to drive spec-aware response generation, and Stage 5 aligns the generator via constrained RLHF. This section describes the framework only; evaluation design is in Section 4, exact settings in the supplementary material, results in Section 5, and the caveats each stage carries are stated with the results they qualify. Table 1 already gives the verdict on each stage, so a reader knows before the description that two of the five do not justify themselves.

3.1. Producing the IR

Stage 1: Intake, Classification, and Aspect Extraction. Stage 1 produces a multi-label class from a RoBERTa head fine-tuned on [1, 15], aspect-level sentiment [15] that keeps “great UI but terrible battery” as two pairs, and entity extraction for devices, OS versions, and named features, which Stage 3 needs as typed fields. The multi-label capacity is almost never exercised: on the 10,000-review generation subset only 5 reviews (0.05%) receive more than one label, and the corpus-scale relabelling artifact stores a single argmax label per review, so the released runs are effectively single-label and readers should not credit the mixed-concern motivation with doing any work in our results. Low-confidence reviews route to HITL-1 by a fixed confidence threshold; letting a model abstain is a familiar idea [8], but we implement a threshold rather than a calibrated selective-prediction procedure and report no risk-coverage behaviour, so the citation is background rather than a description of what runs.

LLM annotation noise is the main challenge. On the 5,038-review praise stratum 25.8% of LLM labels are wrong, and the labeller omits **compatibility** and **performance** almost entirely; per stratum the errors fall in the largest class, not the minority ones, so what a verified anchor repairs is the omission. We apply confident learning [6] against an anchor of 5,230 expert-labeled reviews plus 5,008 from Maalej [1], gated by a confidence threshold on the anchor and a probability ceiling on the LLM label (the supplementary material). Compatibility needs a separate fix, since only 8 of 215,583 LLM labels carry it and Maalej has none: we seed it with 200 template-generated synthetic reviews [33] and 100 keyword-mined

RRGen reviews. The lead author read both sets before training and corrected or dropped items that did not belong, but that pass was done by eye and no correction log was kept, so unlike the other annotation tasks in this paper it leaves no artifact a reader can check. We flag it because 300 of the 310 `compatibility` rows in V5’s training set come from here, so the per-class F1 of 0.83 rests on data whose verification is asserted rather than evidenced.

Stage 2: Three-Layer Knowledge Graph. Stage 2 builds, clusters, and prioritises. The first layer links each review node to its aspects, with edges typed for sentiment and recency where available, but the released run populated less than that: the builder supports entity nodes and signed sentiment edges, yet the run behind our results passes empty entity lists and a constant neutral sentiment, leaving 8,404 review nodes, 10,534 aspect nodes, and no entity nodes. It is in practice the bipartite review-aspect graph PageRank runs over, and the sentiment and entity layers should be read as available in the implementation rather than exercised in the experiments. Aspects come from spaCy noun chunks with a curated keyword list and regular expressions, pruned by frequency and centrality; Hu and Liu [34] is the ancestor of this approach rather than what we implement.

The second layer clusters reviews first by aspect, then sub-clusters within each aspect using sentence embeddings [35] and HDBSCAN [36] directly on the 384-dimensional vectors. Earlier versions of this paper also named UMAP [37] here; that was wrong, since UMAP is used in the flat baseline this design is compared against and not in the released KG-hierarchical run, which applies no dimensionality reduction. Aspect-first keeps each sub-cluster on one theme, giving 605 aspect-grounded sub-clusters of mean size 15.9 against 194 flat clusters of mean size 471.0. An earlier version of this paper reported 3.0 and 51.5; those were the number of representative reviews stored per cluster and a subsample artifact respectively, not cluster sizes. One scale caveat belongs here: the knowledge graph is built over a 10,000-review sample yielding 8,404 usable reviews, so every KG-hierarchical number in this paper describes 3.9% of the corpus, while the flat baseline it is set against clusters 91,368 reviews from the full corpus. The two Stage-2 designs never saw the same input, which is why we do not treat their intrinsic metrics as a head-to-head result (Section 5.4). The third layer prioritises via PageRank [38, 39] on the bipartite review-aspect graph, and each cluster maps to the schema in Table 2 with a centrality-derived priority score.

Stage 3: Review-to-Issue Translation. An LLM agent reads the

standardised cluster schema and writes a specification using the type-routed templates in Table 2: an LLM rather than rule-based extraction because reviews vary too much for fixed rules, and routed by issue type because a bug report needs reproduction steps, a feature request needs acceptance criteria, and a performance complaint needs latency or memory numbers, so a shared template would be incomplete or padded. The five templates are tracker-standard bug reports [5], Cohn user stories [40] with BDD acceptance criteria [41], ISO/IEC 25010 non-functional categories [42], Nielsen heuristics [43], and device-OS matrices [33]. The output passes through HITL-2, where a domain expert scores each spec on a five-dimension rubric (completeness, accuracy, actionability, specificity, clarity, each 1–5), following the LLM-as-judge calibration convention of [44, 45, 46]; five dimensions rather than one, because a single quality number conflates completeness with clarity and the LLM cannot tell which to fix on retry. Any spec failing on any dimension returns for a second pass with dimension-level feedback.

SPECCOV faithfulness metric. We release SPECCOV, an extractive-coverage scorer over substantive-token overlap (length ≥ 5 , stop-word filtered) with quoted-string and coverage bonuses, following news-summarisation conventions [47]; it targets the extractive-coverage sub-type of faithfulness [48], a lexical correlate of factual consistency [49] (other hallucination sub-types in [50] are out of scope). It is a lexical-grounding proxy, not a faithfulness classifier: Section 5.1 reports the experiment that withdraws it as one, and that negative result only makes sense against this definition.

3.2. Using the IR in Response Generation

Stage 4 is a RAG layer over five fixed sources, 15,100 documents in the released ChromaDB index: changelogs and release notes (60 docs), FAQ and help documentation (40), 10,000 past review-response pairs from RRGGen [2], and 5,000 semantically similar past responses, all pre-indexed, plus the validated IssueSpec injected per query. Together they give the generator the answer, the voice, the closest precedent, and the structured intent.

3.3. Aligning the IR-Aware Generator

Stage 5 frames alignment as a CMDP [7, 30],

$$\max_{\theta} \mathbb{E}_{\pi_{\theta}}[R_{\text{quality}}] \quad \text{s.t.} \quad \mathbb{E}_{\pi_{\theta}}[C_{\text{compliance}}] \geq \tau,$$

with the Lagrangian relaxation $\mathcal{L}(\theta, \lambda) = \mathbb{E}_{\pi_\theta}[R + \lambda(C - \tau)]$ collapsing to single-objective optimisation whenever the constraint is already satisfied. We use a CMDP rather than a scalar reward because a dev-rel response can be high-helpfulness and still out of compliance, like a polite promise the team cannot deliver, and a weighted sum hides the trade-off the constraint is meant to expose. The compliance scorer tracks four domain-specific dimensions absent from prior general-safety RLHF [30]: promise safety, information leakage, tone, and legal safety. The trainer is Constrained PPO [29] with a KL-regularised reference policy, with KTO [28] and DPO [27] as small-scale single-objective comparators (the supplementary material). HITL-3 supplies the two feedback streams, quality as reward and compliance as constraint, and also propagates backward, closing the Stage 5→Stage 1 loop in Fig. 1.

4. Experimental Setup

4.1. Dataset, Deduplication, and Annotation Provenance

We use the RRGGen corpus [2]: 310,031 review-response pairs from 58 Android applications. After exact-text deduplication that removes 94,448 duplicate reviews (−30.5%), 215,583 unique reviews form the working corpus, labeled across the seven classes. Evaluation additionally uses a 490-review expert gold standard (70 per class × 7 classes) and 100 stratified clusters for Stage 3.

Of the working corpus, 79.49% of labels come from the V2 LLM annotator with no override, 18.08% are cleanlab-corrected against the verified anchor, and 2.43% are directly human-verified. The lead-author-versus-V2-LLM deviation is 25% on the 5,230-row anchor (1,307 deviations, 24.99%). We previously described this as replicated across two independent strata, the praise stratum and the anchor, on a combined 10,271 rows. That was wrong: the praise stratum is 5,038 of the anchor’s 5,230 rows, so it is 96% of the same sample rather than a second one, and the combined figure double-counted it. The rate is one measurement on one sample, and the praise stratum alone deviates at 25.84% (Table S3 in the supplementary material). This $\approx 25\%$ deviation rate, an order of magnitude higher than i.i.d. label-noise assumptions, motivates the confident-learning procedure of Stage 1.

4.2. Annotation Methodology

Five kinds of annotator contribute, following the multi-annotator convention of [44, 45]. The lead author verified 5,230 reviews across minority

strata, wrote 20 reference IssueSpecs, labeled the 490-review gold blinded to V2 LLM, and scored the Stage 4 sets with method IDs hidden until aggregation; two further human raters annotated the same 490 items and, with her, formed the Stage 4 panel. A Qwen2.5-3B-Instruct judge [46] scored IssueSpecs on the rubric and sub-clusters on the Y/P/N purity rubric, calibrated against the lead author, and the corpus-labeling LLMs (V2, cleanlab-V2, V5) are treated as independent raters for inter-rater statistics. The last two are model raters, not people; three humans contribute in total.

Label provenance, stated plainly. No LLM-generated label is used as ground truth here. LLM labels are the *object under correction*, not the yardstick, and every reported accuracy or κ is measured against a human gold standard; where LLM raters do appear, in the 99-review classification check, they are labelled as such and no headline claim rests on them.

How the gold was constructed, and what that costs. Annotators did not label freely. Each was shown the V5 classifier’s predicted label and asked whether it was correct, supplying a replacement when it was not, under a protocol telling them to accept a reasonable label and overturn only when confident. That makes annotation fast and agreement high, and both flatter us. It also couples the yardstick to the system being measured, since V5’s label is the starting point for every item and V2 and the cleanlab-corrected labels never received that benefit of the doubt. The sample is stratified over V5’s predicted classes as well, forcing 70 items per class into a set where V2 predicts two of those classes almost never, so V2’s per-class F1 of 0.000 on **compatibility** and **performance** is partly a property of the sampling design. The κ trajectory should therefore be read as a comparison on a V5-seeded, class-balanced sample rather than as a neutral benchmark, and a free-labelling gold would settle it. Conflicts are resolved by majority vote: of 489 usable items 476 are unanimous and 13 are two-to-one splits, with no three-way ties, and one item marked wrong by all three raters without a replacement is excluded rather than guessed. All three first completed a shared 20-review calibration round resolved against the written decision rules, and the per-item votes are released.

A disclosure about an earlier annotation round. The 490-review gold was annotated twice. In the first round three sheets were distributed, and the two returned alongside the lead author’s carried annotation columns identical to hers cell for cell, so no genuine agreement could be computed and they are not used anywhere in this paper; the round we report is the second, and the agreement statistics come only from it. We disclose this for

the same reason we disclose the Stage-4 case in Section 5.2: a paper offering a check for non-independent ratings should say where its own artifact failed that check. Both rounds are released.

Stage 4 was likewise rated twice: a first 400-row round whose conditions came from template composers, discarded for the reasons in Section 5.2, and the 200-row round we report, in which both arms come from the same model and differ only in whether the IssueSpec was supplied, all rows scored independently by all three raters blind to condition. Ratings are persisted as spreadsheets with JSON answer keys, blinding shuffles conditions per pair, and the assignment unseals only at aggregation. Table S3 in the supplementary material consolidates agreement across all pairs.

Two three-rater statistics in Table S3 in the supplementary material should not be conflated. The classification gold uses three independent *human* raters and reaches Fleiss $\kappa = 0.968$ on the Y/N verdict and 0.979 on the seven-class label. The $\alpha = 0.285$ row is a separate check on 99 reviews using the lead author plus two Qwen2.5-3B prompts as raters, measuring seven-class classification agreement rather than cluster purity. We reported it as $\alpha = 0.451$ and as a purity check; both were wrong. The value came from an implementation that credited each rating with a coincidence against itself, inflating observed agreement, and the corrected figure agrees with the reference `krippendorff` package to four decimals; the three Cohen’s κ in the same table were computed with scikit-learn and are unaffected. The correction moves the row from the moderate band into the fair band, weakening the LLM-rater panel rather than any claim we make.

4.3. Experimental Design

Because a single-LLM headline raises a structural-bias concern, Stage 3 is replicated across four LLMs from three families, Claude Opus 4.7, Llama-3.3-70B-Instruct (via Groq), and Qwen2.5-{3B, 1.5B}-Instruct (Table S1 in the supplementary material). We pre-register three experiments.

Experiment 1 (RQ1, translation quality). Five conditions on the same 100 stratified clusters: (a) LLM with taxonomy grounding (full Stage 3), (b) free-form LLM, (c) raw concatenation as a lower bound, (d) lead-author reference, and (e) GitHub issues mined from three open-source Android projects as the practitioner upper bound, each scored on the rubric with paired Wilcoxon tests.

Experiment 2 (RQ2), as reported. The design we report holds the generator fixed. The same model, given the same retrieved context,

writes two replies for each of 100 reviews, differing only in whether the validated IssueSpec is in the prompt. All 200 rows are rated blind by three raters on quality, specificity, and helpfulness, and analysed with a paired Wilcoxon test and a paired bootstrap interval. *The discarded design*, described because the paper reports its withdrawal rather than because it supports a claim, used the same 100 clusters to drive four conditions produced by template composers rather than by a model: (a) an rrgen-style baseline (review only), (b) a prompt baseline (CoRe-style dev-rel prompt with lightweight entity extraction), (c) RAG without IssueSpec, and (d) the full system. It was evaluated with automatic metrics (BLEU, ROUGE-L, BERTScore) and 400 blinded paired human ratings, with significance by paired Wilcoxon, Friedman-Nemenyi [51], Bradley-Terry, and McNemar at Bonferroni-corrected $\alpha = 0.05$, and effect sizes following [52, 53]. Both rounds run single-shot vanilla RAG [17] rather than an agentic loop, so the IssueSpec’s contribution is not confounded with orchestration; the agentic comparison against [18, 19] runs on a 10-review feasibility sample and is a demonstration rather than a powered comparison.

Experiment 3 (RQ3, dual- vs. single-objective RLHF). Five policies run on a distilGPT2 proof-of-concept (82M parameters, 400 SFT samples, 296 KTO pairs, 100 DPO pairs): SFT-base, single-objective KTO [28], single-objective DPO [27], a constrained-proxy variant, and dual-objective Lagrangian Constrained PPO [29, 30]. We measure head-to-head BLEU, ROUGE-L, and BERTScore on a 100-review test set, plus the rule-based safety-violation rate.

5. Results

Each research question owns one subsection: Section 5.1 answers RQ1, Section 5.2 answers RQ2 and explains what our earlier positive answer was measuring, and Section 5.3 answers RQ3 in the negative. Sections 5.4 and 5.5 test individual components, and Section 6.1 states which earn their place.

5.1. RQ1: Can a typed IR reach the substantive completeness of human-written GitHub issues?

Under strict content-validity criteria (three or more action-verb reproduction steps, descriptions ≥ 30 words, a full As-a/I-want/So-that triple), the LLM-with-taxonomy condition reaches a template-fill rate of 0.96 (Table 3), free-form LLM 0.69, the same as the lead-author reference set ($n = 20$), raw

concatenation 0.34, and mined GitHub issues 0.53. The LLM-vs-GitHub gap reflects coverage rather than content quality, since the LLM is prompted to fill every field whereas humans leave optional fields blank, and a very-strict sensitivity sweep preserves the ordering, Claude dropping from 0.96 to 0.71 with that gap shifting only -0.085 .

The limit of this measure is worth stating exactly. The free-form condition fills the four shared fields and none of the type-specific ones, so its score is fixed by the issue-type mix, $(30 \cdot 4/7 + 30 \cdot 4/6 + 40 \cdot 4/5)/100 = 0.691$, with zero variance across specs: the 0.69 is a constant, not a measurement, and the gap says only that routing produces type-specific fields at all. Scoring the same conditions on bare field presence gives 1.000 against 0.678, versus 0.959 and 0.678 under the strict criteria, so for the free-form arm the content thresholds remove nothing and the whole reported gap is schema conformance. A handwritten bug specification whose description is “something” repeated thirty-five times, with three five-word steps built from the criterion’s own action-verb list, passes all seven fields and scores 1.000, above our headline. A reader should treat 0.96 as schema conformance and nothing more.

SPECCOV scores, and a correction to how they were computed. An earlier version of our scorer applied per-condition floors, raising `raw_summary` to 5 and the human reference to 4 regardless of content. That is not a measurement, because a score whose value depends on the label of the thing being scored cannot compare labels, so we removed the floors. Across 320 specs the LLM-with-taxonomy condition averages 4.19, free-form LLM 3.38, raw concatenation 4.47, and the human reference 3.45. The two LLM conditions are unchanged, because no floor was ever applied to them; the floors had inflated the two comparison conditions, so removing them makes our own system look better rather than worse, and we say so to put the direction of the correction on the record. On their own these numbers rank raw concatenation highest, the trivial upper bound obtained by copying the source, and Section 5.1 shows they should not be read on their own. Top-ranked aspects by PageRank centrality are `ad` (0.040), `phone` (0.011), and `ads` (0.010), with `edit` (0.009) and `battery` (0.007) next; an earlier version listed `battery` third, skipping the two above it.

SPECCOV does not measure faithfulness

An extractive-coverage score is only useful if it agrees with what a person calls faithful. Three raters independently scored 60 specs, 20 from each of three generators, shuffled and unlabelled, on a 1–5 faithfulness scale, and

separately listed every claim the source reviews do not support. SPEC_{COV} fails all three checks. It does not rank specs the way humans do: Spearman ρ against the three raters is $+0.045$ ($p = 0.73$), $+0.015$ ($p = 0.91$) and $+0.226$ ($p = 0.08$), and $+0.147$ ($p = 0.26$) against their mean. It does not detect the claims humans call unsupported: specs a rater flagged score 3.64 against 3.93 for clean ones, in the right direction but not significant (Mann-Whitney $p = 0.34$). And it orders the generators backwards, ranking the taxonomy condition above the lead-author reference where the humans rank them the other way. We therefore withdraw SPEC_{COV} as a faithfulness measure and release it with this validation attached, so that nobody adopts it on the strength of the construction argument alone. The full analysis, including the per-rater breakdown and the sheet-version qualification that makes us lead with one rater rather than the mean, is in the supplementary material.




Figure 2: Stage-1 classifier recovery and the gold standard it is measured against. *Left*: Cohen’s κ for the three classifier versions, plotted twice, once against the lead-author gold (solid) and once against the three-rater majority-vote gold (hatched). The two are within 0.002 at every step, so the 0.163 to 0.592 trajectory does not depend on any single annotator. Background bands are the Landis-Koch [54] agreement ranges. *Right*: agreement among the three human raters on the same 490 reviews, pairwise and as a three-rater Fleiss κ , all above the almost-perfect threshold. Read the left panel for the result and the right panel for the reliability of the yardstick it uses.

5.2. RQ2: Does conditioning a response generator on the IssueSpec improve human-rated response quality?

An earlier version of this experiment reported that adding the IssueSpec to the response generator raised human-rated quality by 2.35 Likert points.

That number was wrong, not in its arithmetic but in what it measured, and the correction is the most important thing this section has to report.

What was wrong. The four conditions in that experiment were not produced by a language model but by deterministic template composers, Python scripts selecting sentence fragments from a hand-written bank, and those composers differed in authorial register as much as in what information they were given. The `no_spec` arm emitted lemmatised, ungrammatical text averaging 78 words (“we be sorry that the login loop you describe have got in the way”); the `full` arm emitted polished prose averaging 123 words, with cleanup passes and an empathy phrasebook no other arm received. A rater comparing them was separating two writing styles, and the presence of the IssueSpec was confounded with that separation.

The clean comparison, and what we can and cannot verify about it. We re-ran the experiment with the generator held fixed: both arms come from the same model on the same 100 reviews, differing in whether the validated IssueSpec appears in the prompt. One premise of that design is not evidenced by our release and we will not assert it. Neither response file stores the prompt or the retrieved text, and the two record their retrieval sources under different label schemes with different counts, so no released row has an identical recorded source list; both arms draw from the same candidate pool, but whether the retrieved context was identical per review is not recoverable. A reader should treat this as holding the generator and review set fixed rather than as a strict single-variable manipulation. Retrieval differences would add noise rather than a systematic advantage to either arm, which makes a null easier to obtain, so this does not rescue the withdrawn positive result; it does mean the null is weaker evidence than a fully controlled design would give. The two arms average 105 and 97 words with distinct-1 at 0.736 and 0.760, close on the surface features a rater could use to tell them apart, though the length difference is itself significant (paired Wilcoxon $p = 6 \times 10^{-8}$).

The IssueSpec does not measurably improve response quality. Pooled over the three raters the gain is +0.04 Likert points, 95% bootstrap CI $[-0.10, +0.18]$, paired Wilcoxon $p = 0.54$ on 100 pairs, Cliff’s $\delta = 0.04$. The two replies per review are shown in a randomised order; the slot assignment came out 64/36 rather than balanced and raters show a small first-slot preference, so we also computed a slot-stratified estimate, which moves the pooled figure to +0.03 and each rater’s by at most 0.03. We use the unadjusted numbers throughout so that Table 4 is on one estimator, and release the stratified values with the data. Specificity behaves the same way, pooled

-0.00 ($p = 0.87$), and the helpful rate is flat at 94% in both arms. On this design we cannot claim that supplying the typed specification improves the response a user receives, and we should be equally careful in the other direction: with 100 pairs and the observed spread the smallest detectable effect at 80% power is about a fifth of a Likert point, so we have ruled out a large effect and not a small one, and detecting a gain the size of the one we measured would need on the order of 2,500 pairs. The effect, if any, is small enough that this study cannot see it, not zero.

The raters disagree with each other, and that is the interesting part. Table 4 breaks the result out. The lead author, who designed the system, prefers the IssueSpec arm decisively: $+0.56$ quality ($p = 2 \times 10^{-11}$, 59 wins, 5 losses, 36 ties) and $+0.39$ specificity ($p < 10^{-7}$). The two independent raters lean the other way, -0.29 and -0.15 , neither significant at $\alpha = 0.05$ under scipy’s default zero-handling ($p = 0.06$ and 0.09 ; the first has 45 tied rows of 100, and under Pratt’s method that becomes $p = 0.044$, while its bootstrap interval $[-0.56, -0.01]$ excludes zero). We state both rather than pick the one that reads better: on that rater the evidence sits exactly on $\alpha = 0.05$ and moves either side depending on defensible choices. The disagreement between raters is larger than the effect under measurement, and inter-rater reliability is correspondingly low, quadratic-weighted κ [55] of 0.03, 0.04 and 0.50 against 0.97 in the discarded template-based round, where high agreement was a symptom of the confound: when one arm is visibly broken, raters agree easily.

We looked for a content mechanism behind the lead author’s preference and did not find one. Correlating each rater’s per-review preference against the spec-derived features the IssueSpec should introduce, named devices, reproduction steps, and component mentions, produces no correlation that survives inspection; the largest is $\rho = 0.19$ ($p = 0.06$) between the lead author’s preference and reply length. Two readings remain open, that she perceives triage value a general reader has no reason to weigh, in which case the right evaluation asks developers, or that the preference is expectancy on the part of the person who built the system. This design cannot separate them, and we assert neither.

Why we believe the raters this time, and what the check does and does not show. In the discarded round the two additional raters deviated from the lead author on only 16 and 19 of 400 rows, those deviation sets did not overlap on a single row, and every deviation was exactly one point. We initially read the zero overlap as the damning fact. It is not: with

deviations that rare, the hypergeometric probability of zero overlap under independence is 0.45, so it is the single most likely outcome and carries almost no information. The discriminating observation is the second: on the quality column every deviation in both sheets had the same magnitude, one point, which is what a vector perturbed by a fixed step produces and not what human disagreement produces. We name the column because the specificity column in the same sheets does show magnitudes of one, two and three, so the signature is present in one column and not the other. In the reported round the raters differ from the lead author on 133 and 145 of 200 rows, the deviation sets overlap on 101 where chance predicts about 96, and deviations take four distinct magnitudes. We release the check with the data and state its weak prong plainly, because a check offered as a contribution should come with its own limits attached.

What survives. Less than we first thought. The claim that the typed IR carries downstream response quality does not survive, and neither does the retrieval finding we treated as the surviving half. We first reported that retrieval alone was indistinguishable from a no-context baseline, then corrected that to significantly worse; both statements compared the retrieval arm against the *prompt* baseline, which is not a no-context condition because it carries a developer-relations instruction. Against the arm that sees the review text only, retrieval scores 2.24 against 2.31, -0.07 at $p = 0.90$, and all three arms come from the discarded template-composer round in any case, so none supports a claim in either direction. Stage 4 leaves a null and nothing else; Stage 3’s contribution stands on its own evidence in Section 5.1 and no longer has a demonstrated downstream payoff.

5.3. RQ3: Can a CMDP-grounded RLHF loop jointly optimise helpfulness and compliance?

No, and the reason is that the constraint never binds, so the question goes untested rather than answered. Five policies were trained on distilGPT2. Head-to-head on 100 reviews the constrained proxy leads three of four metrics, at +53% BLEU-1 over SFT-base, but that is a single-objective result: at this scale the compliance scorer is satisfied at initialisation ($C = 0.94 \geq \tau$), so λ collapses to zero and the dual update reduces to ordinary optimisation. Tightening τ from 0.50 to 0.90 records 0 binding steps in 30, and a sweep over $\tau \in \{0.0, 0.5, 0.7, 0.9, 1.0\}$ leaves four of five policies at a 100% pass rate. Re-running on a 1.5B model with zero tolerance changes nothing.

A second explanation we cannot separate from the first weakens the finding further. The operational scorer is a regular expression over roughly thirty tokens, and we tested its sensitivity by writing violations by hand. Nine replies that plainly breach the four rubric dimensions, phrased away from the patterns, all score a perfect 1.00, as do empty output and gibberish. What the sweep establishes is that these policies rarely emit about thirty specific tokens; whether they violate the four behaviours the rubric names is not something this instrument can tell us. The compliance figures here are properties of the detector, and the supplementary material carries the full trajectories.

5.4. Cluster quality and purity

Stage 2’s clustering is the one component that loses to its own baseline. A count-matched flat clusterer beats the knowledge graph on all three intrinsic metrics [56, 57, 58] (A1b), and under the same judge the KG’s finer clusters score 0.39 on weighted purity [59, 60] against the flat baseline’s 0.625, though the two audits use different rubrics and sample sizes and are therefore not strictly comparable [61, 62] (supplementary material). We previously reported that the KG achieved $5.4\times$ lower Davies-Bouldin and $1.9\times$ higher Calinski-Harabasz; we withdraw both figures, because the script producing the flat column drew its labels across 194 clusters while embedding reviews from one, so it measured random labelling. The two designs were also never run on the same reviews: the flat baseline clusters 91,368 reviews from the full corpus, the KG covers 5,689 of an 8,404-review sample. What survives is prioritisation, not cluster geometry: the top five aspects by PageRank are touched by 55% of reviews and the top ten by 66%.

5.5. Ablations

An ablation removes one component and re-measures. Eight are listed, counting A1b; seven were run and A7 is deferred. Three came back against the design, and Table 1 carries the verdicts.

A1, remove the knowledge graph and feed flat clusters into Stage 3. Stage 3 quality does not change, because the translator consumes a cluster centroid and its representatives, not the graph edges, so the KG is not a Stage 3 dependency.

A1b, compare the KG against flat clustering matched on cluster count. Intrinsic metrics move with cluster count on their own, so comparing the KG’s 605 sub-clusters against the flat baseline’s 194 conflates structure

with granularity. Re-running agglomerative clustering forced to 605 clusters, on 1,815 representative reviews, removes the confound, and the flat version wins on all three metrics: Davies-Bouldin 1.12 against 2.24 (lower is better), Calinski-Harabasz 4.21 against 1.85 (higher is better), and silhouette +0.148 against -0.234 (a negative value means the average point sits closer to a neighbouring cluster than to its own). The KG does not produce geometrically better clusters; what it produces is an ordering, and prioritisation is the only cluster-layer claim we make.

A2, remove hierarchical sub-clustering. Weighted Y/P/N purity is 0.625 for the flat 194 and 0.390 for the KG’s 605 (Section 5.4): finer granularity costs coherence rather than buying it, the second result against the cluster layer.

A3, remove taxonomy grounding. The LLM writes a free-form issue description instead of filling a routed template. Strict template-fill falls from 0.96 to 0.69 and the per-type structural columns fall to exactly zero (Section 5.1). This is the ablation that isolates the paper’s central mechanism: the gain is not the language model, it is the routing.

A4, remove the HITL-2 spec review. Benchmark scores do not move, for a procedural reason: HITL-2 is a production-time gate and the 100 evaluation specs were scored before any human correction, so the headline numbers already describe pre-HITL output. The ablation says nothing about whether HITL-2 helps in deployment, and we do not claim that it does.

A5, remove retrieval and keep the IssueSpec. BLEU-1 moves by -0.006 , ROUGE-L by -0.008 , and BERTScore F1 by -0.004 , all under one percent absolute. Three caveats belong with this: A5 was generated by the template composer that Section 5.2 withdraws, so it measures retrieval’s effect on a generator we no longer report; it was scored on automatic metrics only while A6 was scored by humans, so the contrast between them is directional and a human-rated no-retrieval condition is the missing cell; and the companion claim that retrieval without the IR falls below a no-context baseline is withdrawn (Section 5.2, -0.07 , $p = 0.90$). A5 therefore reads simply as retrieval making no measurable difference here, with no accompanying story about it doing harm.

A6, remove the IssueSpec and keep retrieval. Superseded. Under the confounded design it showed quality falling from 4.59 to 2.24; under the corrected same-model design (Section 5.2) removing the IssueSpec changes quality by -0.04 , not distinguishable from zero. A5 and A6 together now say something different from what we first concluded: neither retrieval nor

the typed intermediate produces a measurable change in the reply a user receives, and the Stage 3 to Stage 4 coupling we claimed is not supported.

A7, single-stream RLHF feedback. Not run. It tests whether quality and compliance feedback can share one channel, which is only meaningful once the constraint binds and the two signals compete; Section 5.3 reaches the first half of that condition and not the second, so A7 stays deferred rather than reported.

6. Discussion

6.1. What Carries the Result, and What Does Not

Load-bearing. Two components carry the measured gains: type-routed templating, since removing taxonomy grounding (A3) drops strict template-fill from 0.96 to 0.69 and collapses the per-field metrics to zero, and the verified-anchor procedure, whose evidence and caveats are in Section 5.1. The typed IR is *not* a demonstrated driver of Stage 4 quality: removing it under a same-model comparison moves quality by 0.04 ($p = 0.54$, Section 5.2), removing retrieval instead (A5) moves every automatic metric by under one point absolute, and the retrieval-versus-no-context claim we made twice does not survive either. Nothing about retrieval survives from Stage 4 beyond the null.

Not load-bearing, and reported as such. The knowledge graph does not win on intrinsic cluster geometry: under count matching (A1b) flat agglomerative clustering beats it on all three metrics, and sub-cluster purity is lower at the KG’s finer granularity (A2). Its defensible contribution is centrality-based prioritisation, and even that goes less far than it appears, since PageRank runs on the bipartite review-aspect graph and nothing shows it needs the hierarchical sub-clustering layer, a configuration we did not test. Stage 2 as built is therefore justified in prioritisation and unjustified in sub-clustering. Stage 5 sits in between: the constraint does bind once the policy can violate it and the dual update engages, but no quality-versus-compliance trade-off is demonstrated at the scale we can train at (Section 5.3), so it ships as a testbed with a partial result. SPECCOV falls in the same group for a sharper reason: against human faithfulness judgements it shows no correlation at all (Section 5.1), so we withdraw it.

What survives ablation, then, is the typed IR, the type-routed templating that produces it, and the verified-anchor procedure that makes it producible at corpus scale; a reader who wants the smallest system reproducing that

needs Stage 1 and Stage 3. Stage 4 is not part of it, and we state the consequence rather than soften it: on our evidence Stage 4 contributes nothing we can measure, and a reader building on this work should treat the reply-generation layer as unproven rather than as a component with a small effect.

6.2. Implications

Three claims travel beyond app reviews. *First*, LLM annotation noise on software-engineering text is not uniform, but the shape is not the one we first reported. We had described the 25% mislabelling rate as concentrated on minority classes. Recomputing per stratum on the verified anchor shows the opposite: 1,302 of the 1,307 disagreements fall in the **praise** stratum, the largest class, at 25.8%, against 0.5% on **performance**. What is true, and narrower, is that the labeller never emits two of its seven classes at all across the corpus, 8 **compatibility** and 184 **performance** labels in 215,583 reviews. Systematic absence of a class, rather than elevated error on rare classes, is what a small verified anchor is able to repair. *Second*, a caution rather than a claim. We expected a typed structural intermediate to do the load-bearing work in spec-aware response generation and reported it as such; Section 5.2 withdraws that, and the only downstream movement we can still point to is +0.12 on a rule-based scorer from agentic iteration over ten reviews, too small to build on. What generalises is the negative lesson: a structured intermediate can be demonstrably better as an artifact and still leave the generated text indistinguishable, so pipelines of this shape should measure the two separately. *Third*, A1b is a self-contained methodological contribution: whenever a clustering design changes cluster count and structure at once, intrinsic geometry metrics conflate the two, so the minimum reporting standard is a count-matched baseline paired with a domain-utility metric (PageRank coverage here).

6.3. Reproducibility and Extensibility

Every numerical claim in Section 5 is backed by a saved data file, and the full path from raw RRGGen to Stage 4 output takes about six hours on one GPU workstation. We audited this by execution rather than assumption. Re-running the released analysis scripts regenerates roughly sixty stored result files bit-identically, including every statistical test behind RQ1 and RQ2. What those scripts consume is not itself regenerable: the Stage-3 specification sets, the 100-cluster sample and the Stage-4 RAG sample ship with no standalone producer, because producing them needs the proprietary-API

pass. The analysis layer is therefore reproducible from the middle outward and the data-generation layer only against a paid API. One artifact does not pin at all: the GitHub-mined comparison set is re-queried live and returns a different issue mix each run, so the `human_github` condition rests on a snapshot we can release but not regenerate.

The data availability statement itemises the release. We audited the verifier itself, and its coverage is narrower than we previously stated: three segments recompute a value from the underlying data and assert it, while the remainder load a stored summary and print it, which checks that the manuscript agrees with the released file but not that the file is right. Uncovered numbers include the Stage-3 template-fill table, the held-out κ , the purity figures, and most secondary statistics in Sections 5.2 and 5.3.

Deterministic components use fixed seeds; the LLM-dependent stages reproduce statistically rather than bit-identically, and we measured that variance. Regenerating Stage 3 specs three times over the same 15 stratified clusters with the local Qwen2.5-3B generator at temperature 0.7, changing only the seed, gives strict template-fill 0.685, 0.700, and 0.634 (mean 0.673, sd 0.029) and SPECCOV 2.73, 2.67, and 3.00 (mean 2.80, sd 0.144), with no parse failures. Run-to-run noise on template-fill is therefore about three points on a 0–1 scale, and that is the yardstick for the differences we report: the 0.96 versus 0.69 for taxonomy grounding clears it comfortably, the withdrawn Stage-4 effect at +0.04 sits an order of magnitude below it, and gaps of a few hundredths such as 4.09 versus 3.96 in the cross-family rubric do not clear it. We previously also cited the 0.96-versus-0.48 cross-family span here; that comparison is not matched on sample or issue type (Table S1 in the supplementary material), so we no longer use it as a magnitude. Two limits remain: neither the headline Claude generations nor Stage 4 response generation was repeated, because that needs API budget rather than new code, and the released script runs both sweeps unchanged.

6.4. Future work

Five extensions start from released resources. **(0)** Replace SPECCOV with an entailment-based faithfulness measure [63, 64, 65], validated against the 60-spec human ratings we release, and widen the cross-family rubric comparison beyond the 14 matched bug-report clusters with a judge outside all four families. **(1)** Ask the right raters before scaling: a triage reply is consumed by a developer, so the next Stage-4 round should recruit raters with software-engineering experience and no stake in the system before growing

the sample to 200–300 clusters. **(2)** Run the CMDP loop where the trade-off could show, with a longer horizon and a higher base violation rate on an 8B-class model. **(3)** Frontier-family replication (GPT-4o, Gemini-1.5-Pro) on Stages 3–4, and **(4)** cross-corpus transfer of Stages 2–4, V5 already transferring zero-shot to Maalej [1]. **(5)** An agentic planner over Stage 4 dispatching a different sub-agent profile per issue type, using SWE-Agent [20], RepairAgent [21], or HyperAgent [22] for defect repair; a spec-level dispatch stub already runs across all five issue types, and end-to-end patch evaluation is the next step.

7. Threats to Validity

We organise threats using the classification of Wohlin et al. [66], and state after each what we did about it.

7.1. Construct validity

Template-fill measures coverage rather than content correctness, and SPECCOV does not fill that gap (Section 5.1), being a lexical-grounding proxy that does not track human faithfulness judgement [47, 48, 49, 50]. To mitigate this we validated SPECCOV against 60 human faithfulness ratings rather than assuming it worked, and we withdraw it on that evidence rather than reporting it. What remains unmitigated is that Stage 3 then has no validated content-correctness measure at all and the LLM rubric carries that burden alone; NLI- and QA-based measurement [63, 64, 65] is the next step.

7.2. Internal validity

The count-versus-structure confound is exposed and resolved by A1b, which reframes the knowledge graph’s contribution as prioritisation rather than geometry. A threat we did not remove is that the 100 clusters carrying the Stage 3 and Stage 4 results come from the KG pipeline, the layer A1b and A2 find wanting; since flat clustering scores higher on purity, the headline gains may be measured on a slightly worse partition, which would understate rather than inflate them, and re-running Stages 3 and 4 over flat clusters is the cheapest open experiment in the paper. We mitigate this only partially, by reporting the direction of the bias rather than leaving it unstated: because flat clustering scores higher on purity, this threat works against our numbers rather than for them. Two further internal threats are stated with the Stage-1 results: the gold standard is seeded with V5’s predicted label and stratified

over V5’s predicted classes (Section 4), and every `compatibility` figure rests on 300 augmented training rows. We mitigated the first by scoring every classifier against a three-rater majority gold as well as the lead author’s, and the second by auditing the augmentation rather than trusting it, which is how we found the capitalisation cue and withdrew the internal F1 that depended on it (Section 5.1).

7.3. *External validity*

We report no sex- or gender-based analysis. The unit of analysis is a review or a specification, not a person, and the corpora carry no demographic attributes, so no such analysis is possible here. The three annotators are too few for their demographics to support any inference. Whether reviewer or annotator demographics shape the labels or the ratings is an open question this study cannot answer. RRGGen contains 58 Android apps, named by real package identifier rather than anonymised; V5 transfers zero-shot to the Maalej label space, but Stage 2–4 cross-corpus transfer remains unvalidated, GPT-4o and Gemini-1.5-Pro replication is queued (Section 6.4), and every KG-hierarchical number describes a 3.9% sample of the corpus (Section 3.1). We mitigate the single-model threat with a four-model cross-family replication, though its differences are not significant and we report the ordering as descriptive only. The single-corpus threat we do not mitigate; it is the study’s main limit on generality.

7.4. *Conclusion validity*

The three-rater agreement on the classification gold ($\kappa = 0.968$) sits well above what free-labeling studies report, which invites the question of whether the raters were genuinely independent. It also sits in the region where κ becomes sensitive to marginal imbalance [67], so we report raw agreement alongside it. Three properties of the design account for the gap: raters verify a supplied label rather than choose one of seven, collapsing most of the decision space; a 20-review calibration round with adjudicated disagreements precedes the main task; and the protocol fixes eight boundary rules. Residual disagreement is rater-specific rather than random, one rater accounting for most `usability` reassignments, which is the pattern independent annotation produces. We mitigate the remaining doubt by releasing the per-item votes so a reader can recompute any of these figures, and by disclosing that an earlier round of this annotation failed our own independence check and was discarded. Two of the three raters are co-authors, which we cannot mitigate

within this study. *Reliability, and which coefficient belongs to what.* The three agreement statistics here measure different things. The classification gold reaches Fleiss $\kappa = 0.968$ and 0.979 across three human raters. The reported Stage-4 set, rated by the same three humans on all 200 rows, has quality weighted κ of 0.03 , 0.04 and 0.50 ; that is low and it is the number to quote, because the 0.970 to 0.986 range belongs to the discarded round where one arm was visibly broken. The weakest figure, $\alpha = 0.285$ (corrected from a mis-implemented 0.451 , Section 4), comes from a deliberately weaker LLM-assisted panel [44]; no claim rests on it alone. *Statistical.* Experiment 2 reports a null, so the question is power rather than multiplicity: the design rules out a large effect and not a small one (Section 5.2). In the other direction, the one cross-family quality difference reaching nominal significance does not survive multiplicity correction and is withdrawn (Section 5.1).

8. Conclusion and future work

We proposed **IssueSpec**, a knowledge-graph-grounded typed intermediate representation aimed at the five gaps of Section 1, closing the first three and reporting negative results against the last two. A verified-anchor procedure makes the IR producible at corpus scale from noisy LLM labels, lifting Cohen’s κ from 0.163 to 0.592 , and to 0.616 on a held-out subset whose value is replay-dependent and indistinguishable from the full-set figure; classifier endorsement of the corrections is 88.66% , a consistency check rather than independent validation. Cross-family replication on four LLMs preserves the rank order descriptively, with no pairwise difference significant, so it is weak evidence against single-model bias rather than a ruling-out. We set out to show that the typed IR is the load-bearing component of spec-aware response generation, and we cannot: holding the generator fixed moves human-rated quality by $+0.04$ Likert points ($p = 0.54$, three raters, 100 paired reviews) with the raters disagreeing in direction, and the $+2.35$ we previously reported came from a comparison whose arms were written by different template composers. Nor does the negative half of our original hypothesis survive: we had reported that retrieval alone falls below a no-context baseline, and against the arm that actually has no context it scores 2.24 against 2.31 at $p = 0.90$, inside the discarded round in any case. Stage 4 yields one usable result, the null. Two findings travel beyond app reviews: an LLM labeller on software-engineering text can omit whole classes rather than merely err on them, and a structured intermediate can be demonstrably better as an artifact while

leaving the text generated from it indistinguishable, a distinction pipeline evaluations of this shape should measure separately. The released artifacts are listed in Section 6.3.

CRedit authorship contribution statement

Fabiha Jalal: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – original draft, Writing – review and editing, Visualization. **Sadia Tasnim Dhruba:** Investigation, Data curation, Validation, Writing – review and editing. **Tazrian Zaman Tushti:** Investigation, Data curation, Validation, Writing – review and editing. **Ajwad Abrar:** Methodology, Supervision, Writing – review and editing. **Md Kamrul Hasan:** Supervision, Resources, Writing – review and editing. **Hasan Mahmud:** Supervision, Resources, Project administration, Writing – review and editing.

Declaration of competing interest

Two of the three raters of the Stage-4 outputs are co-authors of this paper, and the lead author designed the system she also rated; Section 7 discusses what that does to the result. Beyond that, the authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

All code, configuration files, and analysis scripts are available at <https://github.com/Fabiha-9876/ReviewAgent> [68]. A verification bundle that recomputes or reproduces the numerical claims in Section 5, with the coverage limits stated in Section 6.3, is archived on Zenodo [69] (DOI 10.5281/zenodo.22048738, v3), and the Stage-1 V5 classifier is released on Hugging Face [70]. That bundle supersedes two earlier releases (DOI 10.5281/zenodo.21982774 and DOI 10.5281/zenodo.20320410), the first of which redistributed corpus-scale review text we had no right to redistribute and the second of which also predates the Stage-4 correction described in Section 5.2; both versions remain resolvable, and the earlier one should not be used. The underlying review corpus derives from the RRGGen dataset of Gao et al. That release is access-gated behind a request form, states academic use only, and

carries no licence granting redistribution, so we do not redistribute it. We also do not redistribute the Maalej and Nabil or Guzman corpora. An earlier version of our repository did redistribute all three; we removed them, and `DATA_NOTICE.md` records what was removed and how to obtain each corpus from its original authors. Derived artifacts that embed verbatim review text carry the same restriction as the source and are not offered under an open licence; original code is released under MIT.

Ethics, consent, and data protection

Human participants. Three people annotated for this study: the lead author and two collaborators. The annotation protocol, released with the artifact, states that participation was voluntary, that annotators could withdraw at any time without consequence, that participation was neither required nor rewarded by anything other than the credit their contribution earns, and that no personal information was collected beyond a name and email for attribution. All three are co-authors, and they qualify on contribution: beyond rating, they participated in the calibration round, the adjudication of disagreements, the boundary rules that govern the protocol, and the review of this manuscript, which are the roles the CRediT statement records. We state their dual role explicitly because the paper’s discussion of rater independence must be read alongside it. Annotators were informed of the study’s purpose before starting, and consented to the release of their per-item ratings, which the artifact contains.

The work was not submitted to an institutional review board. Under our institution’s practice a study of this kind, involving colleagues rating text with no collection of sensitive personal data, does not require review. We record that as our own reading rather than as a granted exemption, because no exemption number was issued, and we would rather state the gap than imply an approval we did not obtain.

Review authors. The corpora we work with are user-generated public reviews whose authors did not consent to research use and cannot practically be contacted. We process them under the exemption for scientific research purposes, and took three positions on it. We took three positions on this. We do not redistribute any of the three source corpora, for the licensing reasons in the data-availability statement and because the Maalej and Nabil files contain unredacted email addresses, telephone numbers, and personal names. Derived files we do release contain review text as RRGGen supplies it,

lowercased and lemmatised with digits masked; that masking is imperfect, and we found a small number of surviving transaction identifiers in the source, which anyone re-releasing derived text should re-scan for. We report no result at the level of an individual review author, and no reviewer identifier appears in any released file.

Machine-generated replies. The replies this system produces read as human-written, so any deployment should label them as machine-assisted, which we regard as a condition of use rather than a recommendation.

Role of the funding source

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Acknowledgements

The annotators for this study are among its authors and are credited in the CRediT statement rather than here. The authors verified all technical content, experimental results, and claims in this paper.

Declaration of Generative AI and AI-assisted technologies in the writing process

During the preparation of this work the authors used Anthropic’s Claude [71] in order to assist with manuscript editing, code implementation, and analysis scripting. After using this tool, the authors reviewed and edited the content as needed and take full responsibility for the content of the publication.

An earlier version of this disclosure said Claude assisted only with grammatical refinement and formatting, and that all code and analysis were written by the human authors. The repository contradicts that: the large majority of commits carry a **Co-Authored-By: Claude** trailer, including commits adding pipeline stages, trainers, analysis scripts, and the correction passes of Section 5.2. We correct it here rather than leave a disclosure the artifact refutes. Every structural, experimental-design and claim decision was made by the authors, and every reported number was recomputed from saved data files by the verifier we release.

Generative models are also components of the system under study rather than writing aids, and are reported as such in Section 3 and the supplementary material.

References

- [1] W. Maalej, H. Nabil, Bug report, feature request, or simply praise? on automatically classifying app reviews, in: 2015 IEEE 23rd international requirements engineering conference (RE), IEEE, 2015, pp. 116–125.
- [2] C. Gao, J. Zeng, X. Xia, D. Lo, M. R. Lyu, I. King, Automating app review response generation, in: 2019 34th IEEE/ACM International Conference on Automated Software Engineering (ASE), IEEE, 2019, pp. 163–175.
- [3] J. Dąbrowski, E. Letier, A. Perini, A. Susi, Analysing app reviews for software engineering: a systematic literature review, *Empirical Software Engineering* 27 (2) (2022) 43.
- [4] O. Chaparro, J. Lu, F. Zampetti, L. Moreno, M. Di Penta, A. Marcus, G. Bavota, V. Ng, Detecting missing information in bug descriptions, in: *Proceedings of the 2017 11th joint meeting on foundations of software engineering*, 2017, pp. 396–407.
- [5] T. Zimmermann, R. Premraj, N. Bettenburg, S. Just, A. Schroter, C. Weiss, What makes a good bug report?, *IEEE Transactions on Software Engineering* 36 (5) (2010) 618–643.
- [6] C. Northcutt, L. Jiang, I. Chuang, Confident learning: Estimating uncertainty in dataset labels, *Journal of Artificial Intelligence Research* 70 (2021) 1373–1411.
- [7] E. Altman, *Constrained Markov Decision Processes*, Chapman & Hall/CRC, Boca Raton, FL, 1999, reissued by Routledge, 2021.
- [8] Y. Geifman, R. El-Yaniv, Selective classification for deep neural networks, in: *Advances in Neural Information Processing Systems* 30 (NIPS 2017), 2017.

- [9] G. Bansal, B. Nushi, E. Kamar, W. S. Lasecki, D. S. Weld, E. Horvitz, Beyond accuracy: The role of mental models in human-AI team performance, in: Proceedings of the AAAI Conference on Human Computation and Crowdsourcing, Vol. 7, 2019, pp. 2–11. doi:10.1609/hcomp.v7i1.5285.
- [10] N. Chen, J. Lin, S. C. Hoi, X. Xiao, B. Zhang, Ar-miner: mining informative reviews for developers from mobile app marketplace, in: Proceedings of the 36th international conference on software engineering, 2014, pp. 767–778.
- [11] L. Villarroel, G. Bavota, B. Russo, R. Oliveto, M. Di Penta, Release planning of mobile apps based on user reviews, in: proceedings of the 38th International Conference on Software Engineering, 2016, pp. 14–24.
- [12] A. Di Sorbo, S. Panichella, C. V. Alexandru, C. A. Visaggio, G. Canfora, SURF: Summarizer of user reviews feedback, in: 2017 IEEE/ACM 39th International Conference on Software Engineering Companion (ICSE-C), IEEE, 2017, pp. 55–58. doi:10.1109/ICSE-C.2017.5.
- [13] K. Kaur, P. Kaur, BERT-RCNN: An automatic classification of app reviews using transfer learning based RCNN deep model, Research Square preprint, preprint; not peer reviewed. (2023). doi:10.21203/rs.3.rs-2503700/v1.
- [14] M. A. Hadi, F. H. Fard, Evaluating pre-trained models for user feedback analysis in software engineering: a study on classification of app-reviews, Empirical Software Engineering 28 (4) (2023) 88. doi:10.1007/s10664-023-10314-x.
- [15] E. Guzman, W. Maalej, How do users like this feature? a fine grained sentiment analysis of app reviews, in: 2014 IEEE 22nd International Requirements Engineering Conference (RE), IEEE, 2014, pp. 153–162.
- [16] C. Gao, W. Zhou, X. Xia, D. Lo, Q. Xie, M. R. Lyu, Automating app review response generation based on contextual knowledge, ACM Transactions on Software Engineering and Methodology 31 (1) (2022). doi:10.1145/3464969.
- [17] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel, S. Riedel, D. Kiela,

- Retrieval-augmented generation for knowledge-intensive NLP tasks, in: Advances in Neural Information Processing Systems 33 (NeurIPS 2020), 2020.
- [18] A. Asai, Z. Wu, Y. Wang, A. Sil, H. Hajishirzi, Self-RAG: Learning to retrieve, generate, and critique through self-reflection, in: The Twelfth International Conference on Learning Representations (ICLR), 2024.
 - [19] O. Khattab, K. Santhanam, X. L. Li, D. Hall, P. Liang, C. Potts, M. Zaharia, Demonstrate-search-predict: Composing retrieval and language models for knowledge-intensive nlp, arXiv preprint arXiv:2212.14024 (2022).
 - [20] J. Yang, C. E. Jimenez, A. Wettig, K. Lieret, S. Yao, K. Narasimhan, O. Press, SWE-agent: Agent-computer interfaces enable automated software engineering, in: Advances in Neural Information Processing Systems 37 (NeurIPS 2024), 2024.
 - [21] I. Bouzenia, P. Devanbu, M. Pradel, Repairagent: An autonomous, llm-based agent for program repair, in: 2025 IEEE/ACM 47th International Conference on Software Engineering (ICSE), IEEE, 2025, pp. 2188–2200.
 - [22] H. N. Phan, T. N. Nguyen, P. X. Nguyen, N. D. Bui, HyperAgent: Generalist software engineering agents to solve coding tasks at scale, preprint; not peer reviewed. (2024). arXiv:2409.16299.
 - [23] C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, K. Narasimhan, SWE-bench: Can language models resolve real-world GitHub issues?, in: The Twelfth International Conference on Learning Representations (ICLR), 2024.
 - [24] A. J. W. de Vink, N. Amat-Lefort, L. Han, ReviewGraph: A knowledge graph embedding based framework for review rating prediction with sentiment features, in: 2025 IEEE International Conference on Knowledge Graph (ICKG), IEEE, 2025, pp. 43–50.
 - [25] Z. Yu, S. Wu, J. Jiang, D. Liu, A knowledge-graph based text summarization scheme for mobile edge computing, Journal of Cloud Computing 13 (1) (2024) 9.

- [26] S. Keertipati, B. T. R. Savarimuthu, S. A. Licorish, Approaches for prioritizing feature improvements extracted from app reviews, in: Proceedings of the 20th International Conference on Evaluation and Assessment in Software Engineering (EASE), 2016, pp. 1–6. doi:10.1145/2915970.2916003.
- [27] R. Rafailov, A. Sharma, E. Mitchell, C. D. Manning, S. Ermon, C. Finn, Direct preference optimization: Your language model is secretly a reward model, in: Advances in Neural Information Processing Systems 36 (NeurIPS 2023), 2023.
- [28] K. Ethayarajh, W. Xu, N. Muennighoff, D. Jurafsky, D. Kiela, Model alignment as prospect theoretic optimization, in: Proceedings of the 41st International Conference on Machine Learning (ICML), Vol. 235 of Proceedings of Machine Learning Research, 2024, pp. 12634–12651.
- [29] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, O. Klimov, Proximal policy optimization algorithms (2017). arXiv:1707.06347, doi:10.48550/arXiv.1707.06347.
- [30] J. Dai, X. Pan, R. Sun, J. Ji, X. Xu, M. Liu, Y. Wang, Y. Yang, Safe RLHF: Safe reinforcement learning from human feedback, in: The Twelfth International Conference on Learning Representations (ICLR), 2024.
- [31] M. A. Hearst, Untangling text data mining, in: Proceedings of the 37th Annual meeting of the Association for Computational Linguistics, 1999, pp. 3–10.
- [32] S. Sarawagi, Information extraction, Foundations and Trends in Databases 1 (3) (2008) 261–377. doi:10.1561/19000000003.
- [33] M. E. Joorabchi, A. Mesbah, P. Kruchten, Real challenges in mobile app development, in: 2013 ACM/IEEE International Symposium on Empirical Software Engineering and Measurement, IEEE, 2013, pp. 15–24.
- [34] M. Hu, B. Liu, Mining and summarizing customer reviews, in: Proceedings of the tenth ACM SIGKDD international conference on Knowledge discovery and data mining, 2004, pp. 168–177.

- [35] N. Reimers, I. Gurevych, Sentence-BERT: Sentence embeddings using Siamese BERT-networks, in: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), Association for Computational Linguistics, 2019, pp. 3982–3992. doi:10.18653/v1/D19-1410.
- [36] L. McInnes, J. Healy, S. Astels, hdbscan: Hierarchical density based clustering, *Journal of Open Source Software* 2 (11) (2017) 205. doi:10.21105/joss.00205.
- [37] L. McInnes, J. Healy, J. Melville, UMAP: Uniform manifold approximation and projection for dimension reduction (2018). arXiv:1802.03426, doi:10.48550/arXiv.1802.03426.
- [38] S. Brin, L. Page, The anatomy of a large-scale hypertextual web search engine, *Computer networks and ISDN systems* 30 (1-7) (1998) 107–117.
- [39] L. Page, S. Brin, R. Motwani, T. Winograd, The PageRank citation ranking: Bringing order to the web, Tech. Rep. 1999-66, Stanford InfoLab (1999).
- [40] M. Cohn, *User stories applied: For agile software development*, Addison-Wesley Professional, 2004.
- [41] M. Wynne, A. Hellesøy, S. Tooke, *The Cucumber Book: Behaviour-Driven Development for Testers and Developers*, 2nd Edition, Pragmatic Bookshelf, Raleigh, NC, 2017.
- [42] International Organization for Standardization, ISO/IEC 25010:2023 systems and software engineering — systems and software quality requirements and evaluation (square) — product quality model, Tech. Rep. ISO/IEC 25010:2023, ISO, supersedes ISO/IEC 25010:2011. We cite the 2023 edition because it is the first to enumerate energy and network among the resources covered by resource utilization. (2023).
- [43] J. Nielsen, Enhancing the explanatory power of usability heuristics, in: Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI), ACM, 1994, pp. 152–158, the ten-heuristic set used here. Nielsen and Molich (1990) give an earlier nine-heuristic version; the book *Usability Engineering* (Academic Press, 1993) is a different work.

- [44] F. Gilardi, M. Alizadeh, M. Kubli, Chatgpt outperforms crowd workers for text-annotation tasks, *Proceedings of the National Academy of Sciences* 120 (30) (2023) e2305016120.
- [45] N. Pangakis, S. Wolken, N. Fasching, Automated annotation with generative AI requires validation, preprint; not peer reviewed. (2023). arXiv:2306.00176, doi:10.48550/arXiv.2306.00176.
- [46] L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, I. Stoica, Judging LLM-as-a-judge with MT-bench and chatbot arena, in: *Advances in Neural Information Processing Systems 36 (NeurIPS 2023), Datasets and Benchmarks Track*, 2023.
- [47] M. Grusky, M. Naaman, Y. Artzi, Newsroom: A dataset of 1.3 million summaries with diverse extractive strategies, in: *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, 2018, pp. 708–719.
- [48] J. Maynez, S. Narayan, B. Bohnet, R. McDonald, On faithfulness and factuality in abstractive summarization, in: *Proceedings of the 58th annual meeting of the association for computational linguistics*, 2020, pp. 1906–1919.
- [49] W. Kryściński, B. McCann, C. Xiong, R. Socher, Evaluating the factual consistency of abstractive text summarization, in: *Proceedings of the 2020 conference on empirical methods in natural language processing (EMNLP)*, 2020, pp. 9332–9346.
- [50] Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Bang, A. Madotto, P. Fung, Survey of hallucination in natural language generation, *ACM Computing Surveys* 55 (12) (2023) 1–38. doi:10.1145/3571730.
- [51] J. Demšar, Statistical comparisons of classifiers over multiple data sets, *Journal of Machine learning research* 7 (Jan) (2006) 1–30.
- [52] J. Romano, J. D. Kromrey, J. Coraggio, J. Skowronek, Appropriate statistics for ordinal level data: Should we really be using t-test and

cohen's d for evaluating group differences on the nsse and other surveys, paper presented at the Annual Meeting of the Florida Association of Institutional Research (2006).

- [53] N. Cliff, Dominance statistics: Ordinal analyses to answer ordinal questions, *Psychological Bulletin* 114 (3) (1993) 494–509.
- [54] J. R. Landis, G. G. Koch, The measurement of observer agreement for categorical data, *Biometrics* 33 (1) (1977) 159–174. doi:10.2307/2529310.
- [55] J. Cohen, Weighted kappa: Nominal scale agreement with provision for scaled disagreement or partial credit, *Psychological Bulletin* 70 (4) (1968) 213–220. doi:10.1037/h0026256.
- [56] D. L. Davies, D. W. Bouldin, A cluster separation measure, *IEEE Transactions on Pattern Analysis and Machine Intelligence PAMI-1* (2) (1979) 224–227.
- [57] T. Caliński, J. Harabasz, A dendrite method for cluster analysis, *Communications in Statistics* 3 (1) (1974) 1–27. doi:10.1080/03610927408827101.
- [58] P. J. Rousseeuw, Silhouettes: a graphical aid to the interpretation and validation of cluster analysis, *Journal of computational and applied mathematics* 20 (1987) 53–65.
- [59] E. Amigó, J. Gonzalo, J. Artiles, F. Verdejo, A comparison of extrinsic clustering evaluation metrics based on formal constraints, *Information retrieval* 12 (4) (2009) 461–486.
- [60] A. Strehl, J. Ghosh, Cluster ensembles: A knowledge reuse framework for combining multiple partitions, *Journal of Machine Learning Research* 3 (2002) 583–617.
- [61] C. D. Manning, P. Raghavan, H. Schütze, *Introduction to Information Retrieval*, Cambridge University Press, Cambridge, UK, 2008.
- [62] M. Steinbach, G. Karypis, V. Kumar, A comparison of document clustering techniques, in: *KDD Workshop on Text Mining*, 2000.

- [63] P. Laban, T. Schnabel, P. N. Bennett, M. A. Hearst, Summac: Revisiting nli-based models for inconsistency detection in summarization, *Transactions of the Association for Computational Linguistics* 10 (2022) 163–177.
- [64] O. Honovich, R. Aharoni, J. Herzig, H. Taitelbaum, D. Kukliansy, V. Cohen, T. Scialom, I. Szpektor, A. Hassidim, Y. Matias, True: Re-evaluating factual consistency evaluation, in: *Proceedings of the 2022 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, 2022, pp. 3905–3920.
- [65] E. Durmus, H. He, M. Diab, Feqa: A question answering evaluation framework for faithfulness assessment in abstractive summarization, in: *Proceedings of the 58th annual meeting of the association for computational linguistics*, 2020, pp. 5055–5070.
- [66] C. Wohlin, P. Runeson, M. Höst, M. C. Ohlsson, B. Regnell, A. Wesslén, *Experimentation in Software Engineering*, Springer, Berlin, Heidelberg, 2012. doi:10.1007/978-3-642-29044-2.
- [67] A. R. Feinstein, D. V. Cicchetti, High agreement but low kappa: I. the problems of two paradoxes, *Journal of clinical epidemiology* 43 (6) (1990) 543–549.
- [68] F. Jalal, ReviewAgent: the IssueSpec pipeline, <https://github.com/Fabiha-9876/ReviewAgent>, [software]. MIT licence, code only; see DATA_NOTICE.md. Accessed August 2026. (2026).
- [69] F. Jalal, S. T. Dhruva, T. Z. Tushti, A. Abrar, M. K. Hasan, H. Mahmud, IssueSpec: Reproduction data bundle, version 3, <https://doi.org/10.5281/zenodo.22048738>, [dataset]. Accessed August 2026. (2026). doi:10.5281/zenodo.22048738.
- [70] F. Jalal, IssueSpec V5 classifier, <https://huggingface.co/Fabiha9876/issuespec-v5-classifier>, [software]. Accessed August 2026. (2026).
- [71] Anthropic, Claude (opus 4.7), large language model, <https://www.anthropic.com/claude>, accessed August 2026. Used as a writing and formatting assistant during manuscript preparation. (2026).

Appendix A. Implementation settings and annotator instruments

Every setting behind the results, and both annotator instruments, are in the supplementary material accompanying this submission: the classifier checkpoint and its hyperparameters, the annotator model and prompt behind the noisy corpus labels, the verified-anchor procedure stated as an algorithm, the embedding and clustering parameters, every generator and judge model string with its decoding settings, and every threshold in the scoring instruments. They are also in the released repository. We put them there rather than here because the journal’s length limit does not leave room, and because a reader following the argument does not need them.

Issue type	Template (citation)	Justification
bug_report	title, severity, steps_to_reproduce, expected, actual, component	The tracker-standard field set in GitHub Issues and JIRA. Chaparro et al. [4] support three of the six (observed behaviour, expected behaviour, steps to reproduce). We previously attributed the whole set to Zimmermann et al. [5]; that was wrong twice over. They propose no template, and in their developer survey severity ranks last at 0% importance and component at 3%, against 83% for steps to reproduce. We keep both fields because trackers require them, not because developers report finding them useful
feature_request	John user-story [40] (As-a / I-want / So-that) + behaviour-driven development (BDD) acceptance criteria [41]	Dominant industry format (Agile, Scrum, SAFe); BDD standardises testable outcomes
performance	Fields are latency, startup_time, memory_consumption, battery_drain, network_usage, mapped onto ISO/IEC 25010 [42] performance-efficiency sub-characteristics: the first two to time behaviour, the last three to resource utilisation; capacity is not covered	International standard for software-quality NFR taxonomy. The field names are ours, chosen for mobile-app vocabulary, and are mapped to ISO categories rather than reproducing ISO terms
usability	Nielsen-10 heuristics [43] (visibility, match-real-world, user control, etc.)	Most-cited usability evaluation framework in human-computer interaction
compatibility	Device-OS matrix (devices $\times$ OS versions) [33]	Device/OS fragmentation is the canonical mobile-QA reporting axis

Table 2: IssueSpec standardised issue schema. Each issue type follows the dominant industry or HCI standard. Each template is drawn from a published standard or reference and reproduces its named fields; we ran no conformance check against the standards documents themselves, so “drawn from” should not be read as “certified against”.

condition	strict	bugs:repro	features:story
(a) LLM + taxonomy	0.96	0.73	0.97
(b) LLM free-form	0.69	0.00	0.00
(c) raw concat	0.34	0.00	0.00
(d) lead-author ref ($n = 20$)	0.69	0.00	0.00
(e) human GitHub (3 repos)	0.53	0.13	0.00

Table 3: Experiment 1 substantive template-fill (strict criteria). The 0.96 vs 0.53 LLM-vs-GitHub gap is coverage, not content quality; the strict-vs-very-strict sweep preserves ordering. The exact-zero cells in **bugs:repro** and **features:story** columns for (b), (c), (d), and (e) are structural: free-form LLM, raw concatenation, lead-author, and human-GitHub specs do not emit the template-specific reproduction-step list or As-a/I-want/So-that triple at all when not prompted for them, so the strict-fill rate is zero by construction.

rater	no_spec	full	gain	95% CI	Wilcoxon p
lead author	4.04	4.60	+0.56	[+0.44, +0.68]	2×10^{-11}
independent rater 1	3.87	3.58	-0.29	[-0.56, -0.01]	0.06
independent rater 2	3.57	3.42	-0.15	[-0.33, +0.03]	0.09
pooled	3.83	3.87	+0.04	[-0.10, +0.18]	0.54

Table 4: Stage-4 re-run, same model with and without the IssueSpec, 100 paired reviews, three raters, all 200 rows scored by each. All figures are unadjusted, so the gain column equals the difference of the two condition columns; slot-stratified values are in the released JSON and differ by at most 0.03. The pooled effect is null, and the disagreement between the lead author and the two independent raters is larger than the effect being measured.